

Classification of Chest Diseases using Convolutional Neural Networks

Mellkasi betegségek osztályozása konvolúciós neurális hálózatokkal

Bianka Rác, Kristóf József Varga

Abstract (Hungarian)—Ez a beszámoló a Deep Learning Python és LUA alapokon nevű tárgy keretein belül elkészített nagy házfeladatból hivatott egy átfogó képet adni. A deep learning évről-évre egyre nagyobb teret hódít magának az egészségügy területén, éppen ezért úgy döntöttünk, hogy mi is ezen terület egy problémájának megoldásán szeretnénk dolgozni. A választásunk a mellkasröntgenek elemzésére, a mellkasi betegségek észlelésére esett. Bár az orvoslásban ez hétköznapi feladatnak számít, a gyakorlatban igazán nehézkes lehet. A rendelkezésünkre álló adatokkal és konvolúciós neurális hálók használatával, az álltalunk elkészített program alkalmas megállapítani azt, hogy az adott páciens egészséges e, vagy esetleg további vizsgálatokra szorul.

Abstract (English)—This paper's purpose is to give an overall picture about our project for the course called „Deep Learning Python és Lua alapokon”. The popularity of deep learning in healthcare is increasing year by year. Because of that, we decided to find a solution for one of this field's problems. We wanted to make a program, which can analyze chest x-rays, more precisely to recognize chest diseases. Although, in healthcare it's considered as a common task, in reality it can be really difficult. With the data we have, we can make a neural network which capable to determine if the patient is healthy or further examinations required.



Fig. 1. An image from the National Institutes of Health (NIH) Chest X-Ray dataset, one of the thousands of x-rays we used during our project. The patient has Cardiomegaly.

I. INTRODUCTION

At the beginning of the semester, we weren't sure about what kind of project should we choose. We knew nearly nothing about deep learning and we weren't even sure about what we will be able to accomplish at the end of the course.

The only thing we had back then was the idea that we want to make something in connection with healthcare. First we started to look for databases in the topic. Luckily a lot of people are interested in this field, so there are enormous amount of data about different kind of diseases and also thousands of studies and contemporary histories. Soon we found an interesting bone x-ray dataset, but unfortunately it was private and really hard to access.

Finally we chose the dataset of the National Institutes of Health (NIH) about Chest X-Ray images. Our goal was to make the model be able to predict whether the patient has any kind of lung disease and if has, what type.

Before we could do anything in this project, we had to learn about convolutional neural networks and other topics of deep learning. Of course it wasn't easy at first, but our motivation didn't decrease. After a few weeks we started to inspect the dataset even more, and tried to find out what should we do with it, before the implementation of the neural network.

The preprocessing was a crucial part of the project. We generated different kind of .csv files, in order to test out different scenarios and to identify which attributes should we use.

Our final goal was clear, but we wanted to start with something easier first. So we decided to implement a network which can predict if the patient has something abnormal in the chest, or not. (Healthy or not.)

II. RECENT EFFORTS

There have been recent efforts on creating projects for this cause, with different complexity levels. The biggest one, is made by the same research group who made the NIH dataset (Xiaosong Wang, Yifan Peng, Le Lu, Zhiyong Lu, Mohammadhadi Bagheri, Ronald M. Summers). It's a really successful and well planned project with amazing results, but sadly it's not openly available.

We found other implementations which were public. Of course we wanted to make our own implementation, but they were a good starting point to find and recognise the possible problems with the dataset.

III. DATASET

There are different kind of datasets related with chest x-rays, but most of them are old, poorly managed or includes low quality images. As we mentioned before, our final choice was the chest x-ray dataset by NIH.

We found the dataset on Kaggle, and it was really easy to get it for free (we also discovered, it's way easier to press the "Download" button on Kaggle, than trying to download the dataset from code, or from the terminal). The size of the dataset is 42 GB and there are 112,120 images in it, with the size of 1024 x 1024. The dataset consists of scans from more than 30,000 patients.

The labels were made by the radiologist of NIH Clinical Center. The labels are in a csv file (Data_entry_2017.csv) with some added features like age, number of follow up visits, AP vs PA scan, and the patient's gender.

The structure of the Data_entry_2017.csv file:

- Image Index: File name
- Finding Labels: Disease type (Class label)
- Follow-up #
- Patient ID
- Patient Age
- Patient Gender
- View Position: X-ray orientation
- OriginalImageWidth
- OriginalImageHeight
- OriginalImagePixelSpacing_x
- OriginalImagePixelSpacing_y

We didn't use every one of these attributes. The most important for us were the Image Index, which basically was the name of the file, and the Finding Labels, which showed us what kind of problem can be seen on the picture.

In the dataset, there are 14 types of diseases, so we have 15 classes (with No Finding):

- Atelectasis
- Consolidation
- Infiltration
- Pneumothorax
- Edema
- Emphysema
- Fibrosis
- Effusion
- Pneumonia
- Pleural_thickening
- Cardiomegaly
- Nodule
- Mass
- Hernia
- No Finding

During the implementation of the first stage, we only wanted to do a binary classification, so in this case we used all of the diseases names as one category.

IV. PREPROCESSING

We tried numerous way to train our model, but the enormous amount of the data caused some problem to us. We had about 110,000 images in our dataset, it was very slow to load or upload them, and our devices didn't have enough capacity to them. Our final decision was that we don't use the whole dataset for the first stage, just a piece of it – about 10,000 images.

We started the preprocessing with the Data_entry_2017.csv file. We renamed it to Data_Entry.csv. In this dataset we can find the names of the images, the category of the illness, and some information about the patients. We edited the column names by hand because they contained spaces, but later we used preproCsv.py Python file. The owners of the linked gitHub repository mentioned that there are some corrupted data in the dataset: some ages were in days or months, so we deleted them, apprx. 27 rows. We deleted some columns (Original_Image_Width, Original_Image_Height, Original_Image_Pixel_Spacing_X, Original_Image_Pixel_Spacing_Y) because they don't have any effect on data.

In the original file there were apprx. 709 different categories, because some images belonged to more than one class. We deleted these rows to reduce the number of categories, so 20,729 rows were deleted. There were some rows with unrealistic age (more than 120 years), so 16 rows became deleted.

No Finding	60387
Infiltration	9544
Atelectasis	4210
Effusion	3959
Nodule	2706
Pneumothorax	2198
Mass	2137
Consolidation	1314
Pleural_Thickening	1127
Cardiomegaly	1094
Emphysema	895
Fibrosis	727
Edema	633
Pneumonia	307
Hernia	110

Fig. 2. The number of images in each categories after reducing the categories.

We deleted the corrupted data so we had to choose the 10,000 rows of image, which we want to use to train and test our model. We had two classes: healthy or not. First we chose 5,000 rows about images of healthy lungs randomly, and another 5,000 rows about images of lungs with disease. We added two other columns to the data frame with labels positive and negative. Positive column indicates lungs with disease, and negative indicates healthy lungs. We used one-hot encoding for encoding these categories. After that we shuffled the whole dataset and separated it into three different dataset: train.csv – this is the training dataset, valid.csv – this is the validation dataset and test.csv – this is the testing dataset. We used three different dataset because we don't need to work with the whole dataset during the teaching.

The next step was the preprocessing of the images. We wrote a script named prepro.py. This script's main responsibilities are image preprocessing and image selection. First, we had to load the preprocessed CSV file – entry_data_edited.csv – into a data frame. In this way, we could easily iterate through the whole dataset, row by row. Because the rows of the CSV file had been already mixed, we could separate the files into folders - train, valid, test – one after another, without any kind of randomization. This means, we could put the first 50% of data into the train folder, then the next 20% of the data into the valid folder, and the remaining 30% of the data into the test set. .

Before we put our pictures to one of the folders, we had to preprocess them. For image loading, we used openCV. We loaded them as grayscale images, than we resized them to 256px * 256px.

V. THE MODEL

We used Keras with Tensorflow backend. For our project we created a Convolutional Neural Network (CNN). It has four convolutional, two max-pool and two fully-connected layers:

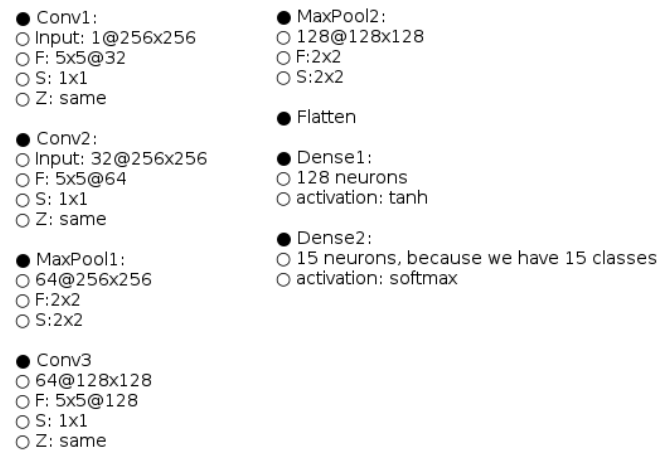


Fig. 3. The first version of our modell.

It's important to point out, that this network is just one of the networks we tried out. It's for a „full” classification with 15 output classes, this is for our final goal. But because of the lack of time and resources we had to work on smaller network, or with fewer outputs.

We tried out different kind of activation functions, and even built a network to solve this problem as a regression.

VI. RESULTS

We tried out many versions of our model. Unfortunately we couldn't try out with the whole original dataset with 15 classes because our devices didn't have enough capacity. First we used just a small part of the dataset (3,7%).




Fig. 4. The measure of loss of the first version

As you can see the The model didn't learn well enough: the mean-squared error was 0.037786, the validation loss was 1.30943, and the accuracy was 0.6652. There are more reasons why. First, we used just a little part of the whole dataset (3,7%). It's possible that the model didn't get sample of every classes during the training phase, but during the validation or test phase did.

In the next step we worked with just 5,000 from four classes: one healthy and three disease. We didn't get good results, so we started to deal with hyperparameter optimization. First we added L1 regularization, but it didn't helped us, so we didn't used it later. After that we added Dropout(0.5) layer after the first dense layer.




Fig. 5. The measure of accuracy of the second version




Fig. 6. The measure of loss of the second version

As you can see at Fig y. and x. the validation loss and accuracy improved, it was about 0.5 and 1.2. But it should have stopped about at the 30th epoch, because it didn't improved later.

After that we reduced the number of epochs to two (healthy or not).




Fig. 7. A




Fig. 7 B. It improved a lot the validation loss was 0.664 and the validation accuracy 0.6105

VII. WHAT'S NEXT

For us this project was more about experimenting, than building a world-class solution for a problem. We think this is a really interesting project, an what we were able to achieve in this semester is just the beginning. We definitely want to continue what we started, hopefully with better results.

In the future, we should train out network way longer. It's a quite complex task and we saw in case of some of the network, there is a slow but sure learning curve, which is recognizable even at the beginning of the training process.

For better results, we could use the other attributes of the dataset, like the age and gender of the patients.

VIII. REFERENCES

- [1] Mire használják a deep learninget, és a CNN-t is az orvostudományban
<https://www.annualreviews.org/doi/pdf/10.1146/annurev-bioeng-071516-044442>
- [2] 3féle cnn-t próbáltak ki, lung disease classification-ra is
<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC4890616/>
- [3] Dropout az overfitting ellen
http://www.jmlr.org/papers/volume15/srivastava14a/srivastava14a.pdf?utm_content=buffer79b43&utm_medium=social&utm_source=twitter.com&utm_campaign=buffer
- [4] Momentum & Nesterov momentum
<http://proceedings.mlr.press/v28/sutskever13.pdf>
- [5] A mi database-ünkkel foglalkozik, de multi-label classificationt használt és mi meg nem
http://openaccess.thecvf.com/content_cvpr_2017/papers/Wang_ChestX-ray8_Hospital-Scale_Chest_CVPR_2017_paper.pdf
- [6] Dataset leírás
<https://www.kaggle.com/nih-chest-xrays/data/home>
- [7] Neurális hálózatok – Altrichter Márta, Horváth Gábor, Pataki Béla
https://www.tankonyvtar.hu/hu/tartalom/tamop425/0026_neuralis_4_4/0026_neuralis_4_4.pdf
- [8] Deng L.
<http://research.microsoft.com/pubs/209355/DeepLearning-NowPublishing-Vol7-SIG-039.pdf>